

Predictive Modeling in Multiclass Classification

Steven Futter

2/13/2017

In this paper I fit a model suite of the following models: Random Forest, a Support Vector Machine, and a Neural Network model. The data is the 'Wine' data set that can be downloaded from the [Irvine Machine Learning Repository](#). Since this is a small data set I only focus on fitting models in-sample.

(1) Data Quality Check

The wine data set consists of 178 observations and 14 variables. Of the 14 variables there is one class variable and 13 continuous variables that describe the profile of the wine. In this assignment we are dealing with a classification problem since the goal of the assignment is to predict which class a wine falls into based upon the continuous variables, such as alcohol level, magnesium content, and color intensity.

Missing Data, Data Ranges and Distributions

Of the 178 observations the larger portion of wines belong to class 2 (40%), followed by class 1 (33%), and then class 3 (27%). The `my.summary()` function below confirms that there are no missing observations in the wine data. Based upon the mean and median values the proline variable appears to have a non-normal distribution since mean and median values differ relative to other variables in the data set. The mean and median values of the proline variable are noticeable different. The mean value of proline is 746.89 whereas the median value is 673.50. Other variables in the data set appear to have extreme observations appearing as minimum and maximum values. There are no negative minimum values which makes sense and although there are outliers the values appear to be valid at first glance.




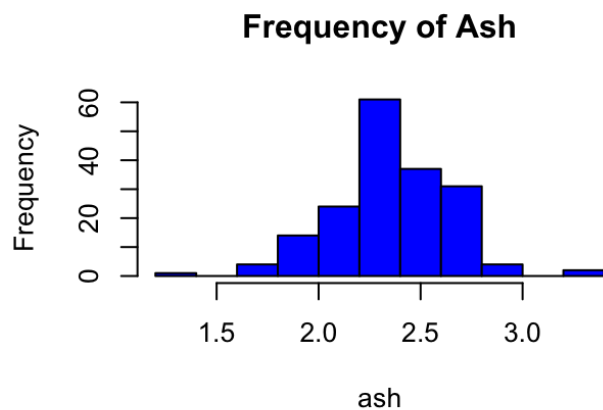
```
my.summary(wineData[,c(2:14)])
```

```
##          alcohol malicAcid  ash alkalinityOfAsh magnesium totalPhenols
## max          14.83      5.80 3.23              30.00    162.00      3.88
## min          11.03      0.74 1.36              10.60     70.00      0.98
## 0.01         11.44      0.90 1.70              11.35     78.00      1.14
## 0.05         11.66      1.06 1.92              14.77     80.85      1.38
## 0.25         12.36      1.60 2.21              17.20     88.00      1.74
## 0.5          13.05      1.87 2.36              19.50     98.00      2.36
## 0.75         13.68      3.08 2.56              21.50    107.00      2.80
## 0.95         14.22      4.46 2.74              25.00    124.30      3.27
## 0.99         14.47      5.54 2.99              28.50    141.76      3.60
## mean         13.00      2.34 2.37              19.49     99.74      2.30
## variance      0.66      1.25 0.08              11.15    203.99      0.39
## % missing     0.00      0.00 0.00              0.00      0.00      0.00
##          flavanoids nonflavanoidPhenols proanthocyanins colorIntensity
## max              5.08                  0.66              3.58          13.00
## min              0.34                  0.13              0.41          1.28
## 0.01             0.47                  0.14              0.42          1.86
## 0.05             0.55                  0.19              0.73          2.11
## 0.25             1.20                  0.27              1.25          3.22
## 0.5              2.13                  0.34              1.56          4.69
## 0.75             2.88                  0.44              1.95          6.20
## 0.95             3.50                  0.60              2.71          9.60
## 0.99             3.79                  0.63              3.03         11.02
## mean            2.03                  0.36              1.59          5.06
## variance         1.00                  0.02              0.33          5.37
## % missing        0.00                  0.00              0.00          0.00
##          hue OD280_OD315OfDilutedWines proline
## max          1.71                  4.00 1680.00
## min          0.48                  1.27 278.00
## 0.01         0.55                  1.29 306.94
## 0.05         0.57                  1.46 354.55
## 0.25         0.78                  1.94 500.50
## 0.5          0.96                  2.78 673.50
## 0.75         1.12                  3.17 985.00
## 0.95         1.28                  3.58 1297.25
## 0.99         1.43                  3.84 1522.36
## mean         0.96                  2.61 746.89
## variance      0.05                  0.50 99166.72
## % missing     0.00                  0.00 0.00
```

By plotting a matrix of histograms we can visually identify clear outliers in the data set that are described above by the `my.summary()` function. Outliers can be seen in the following variables: ash, magnesium, flavanoids, hue.

Outliers: i. left tail of ash ii. right tail of ash iii. right tail of magnesium iv. right tail of flavanoids v. right tail of hue


Outliers



(2) Exploratory Data Analysis

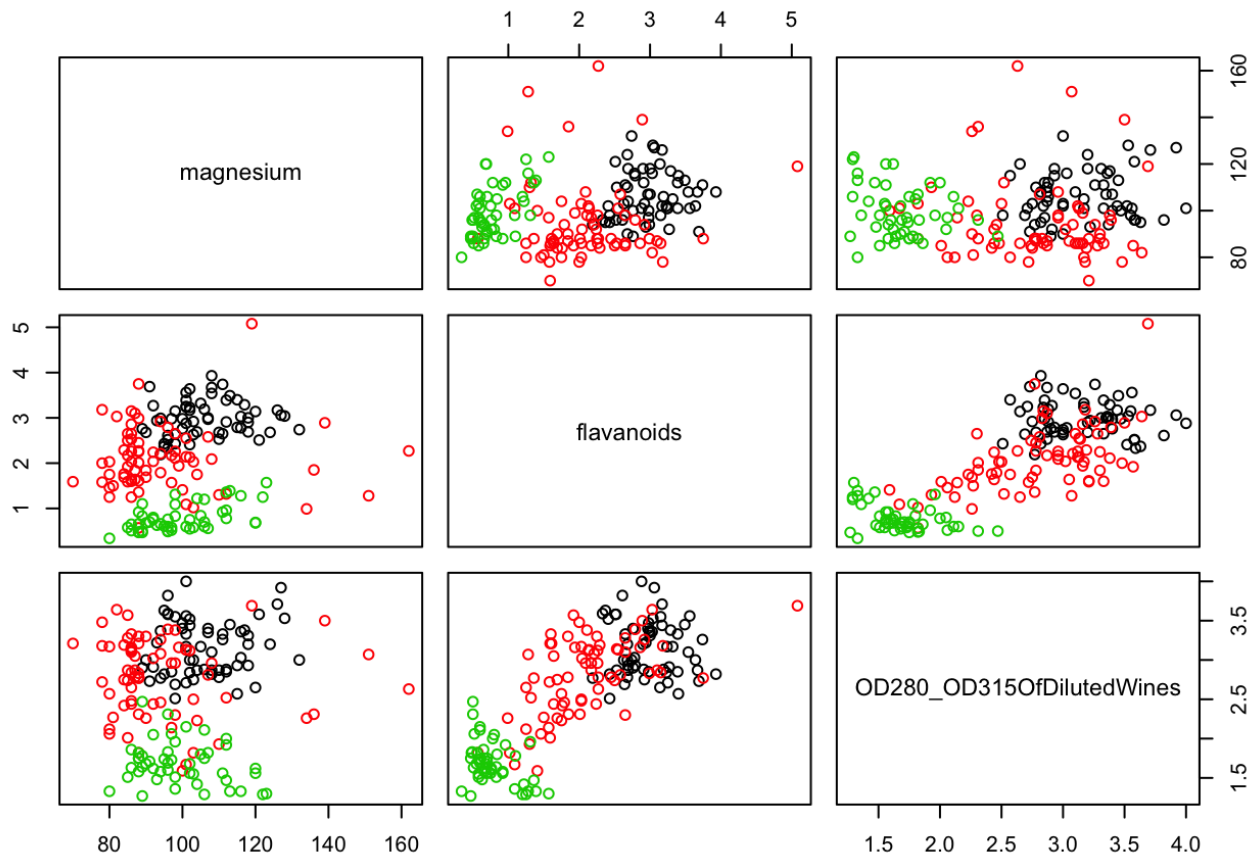
With the wine data set we have a classification problem which can benefit from viewing the data distributions with boxplots. From the boxplots below we can see that alcohol, totalPhenols, and flavanoids appear to be particularly predictive with the middle 75% (grey shaded bar) of values separating the three wine classes with minimal overlap. This looks promising since the wine classes are differentiated quite nicely by these three variables alone.

Boxplots: Alcohol, totalPhenols, and flavanoids differentiated across wine classes



Several of the variables appear to be good candidates for a discriminative classifier such as a Support Vector Machine (SVM). An SVM uses an algorithm to split the data classes into optimal hyperplanes. Below I show that various pairs of variables cluster the wine classes into more or less separate wine classes. From initial EDA it appears that a clustered algorithm approach to modeling may also have a low misclassification error percentage. Below we can see how magnesium, flavanoids, and OD280_OD315OfDilutedWines pairs segregate the three wine classes into separate clusters.

Scatterplots: magnesium, flavanoids, and OD280_OD315OfDilutedWines pairs segregate the wine classes




Using the Lattice library we can see the density function of each wine class against alcohol. We can see that the peak densities of each wine class vary as alcohol concentration increases. Generally speaking we can see that class 2 wines have the least alcohol content and class 1 wines have the most.

Density Plot Alcohol by Class

Wine Class


1 — 2 — 3 —



I now use a tree model to explore the data further and evaluate missclassification error.

Tree Model

```
##  
## Classification tree:  
## tree(formula = class ~ ., data = wineData)  
## Variables actually used in tree construction:  
## [1] "flavanoids"      "colorIntensity" "malicAcid"      "proline"  
## [5] "alcohol"  
## Number of terminal nodes: 7  
## Residual mean deviance: 0.08779 = 15.01 / 171  
## Misclassification error rate: 0.01685 = 3 / 178
```



The tree model above has a misclassification of 1.7%. flavanoids, proline, colorIntensity, malicAcid, and alcohol are important variables in determining the wine's class. Wine class '3' is segregated along the left branch with flavanoid values less than 1.575 and color intensity greater than 3.825. Class 1 and 2 are split either side of the first tree root.

Fitted Model Suite:

1. Random Forest

Firstly, let's create a random forest model taking in all variables within the wine data set. I write them out in the model to make it explicit which variables are being included into the random forest model.

```

Call:
  randomForest(formula = class ~ alcohol + malicAcid + ash + alcalinityOfAsh +
    magnesium + totalPhenols + flavanoids + nonflavanoidPhenols +      proanthocyanins
    + colorIntensity + hue + OD280_OD315OfDilutedWines +      proline, data = wineData,
    mtry = 2, importance = TRUE)

      Type of random forest: classification
      Number of trees: 500
No. of variables tried at each split: 2

      OOB estimate of  error rate: 1.12%
Confusion matrix:
   1  2  3 class.error
1 59  0  0 0.00000000
2  1 69  1 0.02816901
3  0  0 48 0.00000000

```

0% missclassification can be seen in the confusion matrix table.

```
randomForest.pred  1  2  3
                  1 59  0  0
                  2  0 71  0
                  3  0  0 48
```

2. Support Vector Machine

Let's now run a SVM model and evaluate the confusion matrix below. Similarly to the Random Forest model the SVM produces a 0% misclassification error. See the confusion matrix table below.

```
Call:
svm.default(x = x, y = y, data = wineData)
```

```
Parameters:
  SVM-Type:  C-classification
SVM-Kernel:  radial
    cost:    1
   gamma:    0.07692308
```

```
Number of Support Vectors:  69
```

```
( 19 31 19 )
```

```
Number of Classes:  3
```

```
Levels:
 1 2 3
```

```
      y
svm.pred  1  2  3
         1 59  0  0
         2  0 71  0
         3  0  0 48
```

3. Neural Network Model

Finally, I run a neural network model using the nnet library. Again, I produce a confusion matrix to evaluate the model's performance.


```
# weights:  428
initial  value 213.223719
iter   10 value 136.377871
iter   20 value  95.624574
iter   30 value  26.037956
iter   40 value  21.125682
iter   50 value  20.389365
iter   60 value  19.173270
iter   70 value  15.030089
iter   80 value  12.598054
iter   90 value  11.187454
iter  100 value  11.138673
final   value  11.138673
stopped after 100 iterations
```

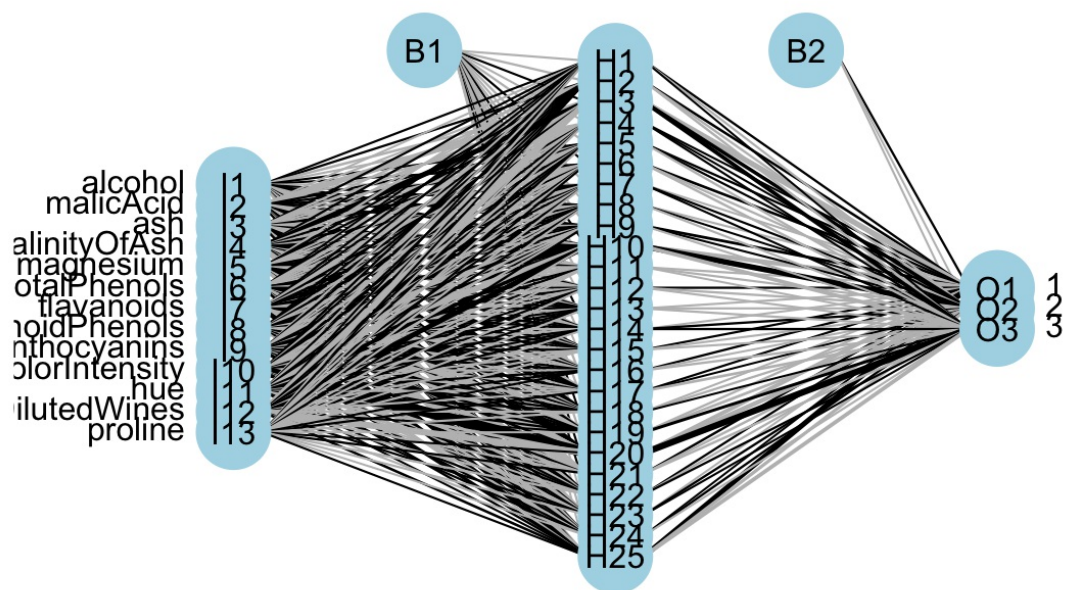
```
pred
  1  2  3
1 57  2  0
2  0 70  1
3  0  1 47
```

```
[1] 0.9775281
```

A plot of a neural network can be produced using the `plot.nnet` function, however, it's output is best to demonstrate that 25 nodes were created in the process. These nodes can be seen in the middle of the plot below.

Neural Network Plot: 1 layer separating variable inputs and class outputs

The plot below highlights that different weights are added to each model input to produce the neural network model's class output prediction.



Results

Since each model is a 'black box' when it comes to determining how the inputs are treated, it is important that each model is compared for predictive accuracy. To compare each model I present the output table below which provides the misclassification error rate for each model. As the assignment asked to evaluate results in-sample I provide the misclassification error rate for each model using the entire set of wine data observations.

	model	model.errors
1	Random Forest	0.00000000
2	Support Vector Machine	0.00000000
3	Neural Network	0.02247191

As can be seen from the results table the Random Forest and SVM models have excellent results — both have a 0% misclassification error. The Neural Network model did not perform as well, but its misclassification error was still very low with. Misclassification error for the Neural Network model is 2.25%.

Appendix

```

# 1. Random Forest
set.seed(1)
randomForest.fit = randomForest(class~
  alcohol + malicAcid + ash + alcalinityOfAsh +
  magnesium + totalPhenols + flavanoids + nonflavanoidPhenols + proanthocyanins
+
  colorIntensity + hue + OD280_OD315OfDilutedWines + proline,
  data=wineData, mtry=2, importance=TRUE)
randomForest.fit

# 2. SVM
x = subset(wineData, select=-class)    # Divide wineData to x (the variables) and y
the classes
y = wineData$class

svm.fit = svm(x,y,data=wineData)      # Create SVM model
summary(svm.fit)

svm.pred=predict(svm.fit, x)
table(svm.pred, y)
svm.train.error = 1 - mean(svm.pred==y)
svm.train.error

# 3. Neural Net
#The neural network requires that the species be normalized using one-of-n normaliz
ation. We will normalize between 0 and 1. This can be done with the following comma
nd.
ideal <- class.ind(wineData$class)

# Train the neural network
set.seed(123)
wineANN = nnet(wineData[,-1], ideal, size=25, softmax=TRUE)

# test the output from the neural network on the validation data
pred = predict(wineANN, wineData[,-1], type="class")
table(wineData[,1], pred)

# Plot tutorial found here: https://beckmw.wordpress.com/2013/11/14/visualizing-neu
ral-networks-in-r-update/
# Plot the neural net output
install.packages('devtools')
install.packages('reshape')
library(devtools)
require(reshape)
source_url('https://gist.githubusercontent.com/fawda123/7471137/raw/466c1474d0a505f
f044412703516c34f1a4684a5/nnet_plot_update.r')

#plot the model
plot.nnet(wineANN)

```